

# DelayedMatrixStats

Porting the matrixStats API to work with  
DelayedMatrix objects

Peter Hickey

2017-07-26

# Why matrixStats?

**matrixStats** by Henrik Bengtsson and co. on CRAN since 2009

“Functions that Apply to Rows and Columns of Matrices”

## Optimised row/column operations on *matrix* objects

```
# Simulate some zero-inflated count data
matrix <- matrix(sample(0:100, 20000 * 100, replace = TRUE),
                 nrow = 200,
                 ncol = 10000)
matrix[sample(length(matrix), length(matrix) * 0.6)] <- 0L

# matrixStats has fast, memory-efficient column summaries
benchmark(apply(matrix, 2, median),
          matrixStats::colMedians(matrix),
          times = 1)

#>                               expr Median time (ms) Mem alloc (MB)
#>      apply(matrix, 2, median)          774.0          50.2000
#> matrixStats::colMedians(matrix)         28.7           0.0825
```

2 / 12

# Why **matrixStats**?

## Optimised row/column operations on *matrix* objects

```
library(matrixStats)

# matrixStats optimised for subsetting calculations
j <- c(2001:3000, 5001:5500)
benchmark(colSums(matrix[, j]),
          colSums2(matrix, cols = j),
          times = 1)

#>               expr Median time (ms) Mem alloc (MB)
#>   colSums(matrix[, j])           5.22         1.2100
#> colSums2(matrix, cols = j)        1.19         0.0171
```

# Why matrixStats?

## Lots of useful col/row summary functions

```
grep("^col", getNamespaceExports("matrixStats"), value = TRUE)
#> [1] "colMadDiffs"      "colCummins"      "colRanks"
#> [4] "colWeightedVars"  "colQuantiles"    "colDiffs"
#> [7] "colCumprods"      "colSds"          "colCollapse"
#> [10] "colVars"          "colAnyMissings"  "colWeightedSds"
#> [13] "colCummaxs"       "colAlls"         "colVarDiffs"
#> [16] "colIQRs"          "colMins"         "colWeightedMedians"
#> [19] "colLogSumExps"    "colAvgPerRowSet" "colSdDiffs"
#> [22] "colIQRDiffs"      "colSums2"        "colCumsums"
#> [25] "colTabulates"     "colMedians"      "colOrderStats"
#> [28] "colWeightedMads"  "colMaxs"         "colCounts"
#> [31] "colWeightedMeans" "colMeans2"       "colProds"
#> [34] "colRanges"        "colAnyNAs"       "colAnys"
#> [37] "colMads"
```

# Big data blues

- You've got matrix-like data but too large for in-memory *matrix* :(

# Big data blues

- You've got matrix-like data but too large for in-memory *matrix* :(

## *DelayedMatrix!*

- A wrapper around a matrix-like object
- Data can be in memory or on disk
- *DelayedMatrix* works as an assay in a *SummarizedExperiment*
- *DelayedMatrix* supports the standard & familiar *matrix* API<sup>\*</sup>
  - [
  - dim()
  - dimnames()
  - t()
  - log()
  - **colSums()**
  - ...

[\*] But not subassignment

# *DelayedMatrix* backends

## In-memory backends

```
DelayedMatrix <- DelayedArray::DelayedArray(matrix)
pryr::object_size(DelayedMatrix)
#> 8 MB

DelayedddgCMatrix <- DelayedArray(as(matrix, "dgCMatrix"))
pryr::object_size(DelayedddgCMatrix) # Larger than dense version!
#> 9.55 MB

RleMatrix <- RleArray(Rle(matrix), dim = dim(matrix))
pryr::object_size(RleMatrix)          # Low RLE compressibility
#> 10.1 MB

TricksyRleMatrix <- as(matrix, "RleMatrix") # Uses tricky tricks
pryr::object_size(TricksyRleMatrix) # Tricky tricks in play
#> 6.34 MB
```

# *DelayedMatrix* backends

## On-disk backends

```
HDF5Matrix <- HDF5Array::writeHDF5Array(matrix)
pryr::object_size(HDF5Matrix)
#> 2.39 kB
file_size(HDF5Matrix@seed@file)
#> 1.63 MB

matterMatrix <- matterArray::writeMatterArray(matrix)
pryr::object_size(matterMatrix)
#> 9.63 kB
file_size(matterMatrix@seed@matter@paths)
#> 8 MB
```



# Why DelayedMatrixStats?

# Why DelayedMatrixStats?



# Why **DelayedMatrixStats**?

- Support **matrixStats** API for *DelayedMatrix* and derived classes
- Reduce friction between using *matrix* or *DelayedMatrix*

# Why **DelayedMatrixStats**?

- Support **matrixStats** API for *DelayedMatrix* and derived classes
- Reduce friction between using *matrix* or *DelayedMatrix*

## Initial release aim

General 'block-processing' method to work for *DelayedMatrix* and arbitrary derived classes

# Why **DelayedMatrixStats**?

- Support **matrixStats** API for *DelayedMatrix* and derived classes
- Reduce friction between using *matrix* or *DelayedMatrix*

## Initial release aim

General 'block-processing' method to work for *DelayedMatrix* and arbitrary derived classes

## Subsequent releases

'Backend-aware' optimised methods

# Why DelayedMatrixStats?

Yay, same syntax works regardless of backend!

```
benchmark(colMedians(matrix),
          colMedians(DelayedMatrix),
          colMedians(DelayedddgCMatrix),
          colMedians(RleMatrix),
          colMedians(TricksyRleMatrix),
          colMedians(HDF5Matrix),
          colMedians(matterMatrix),
          times = 1)

#>           expr Median time (ms) Mem alloc (MB)
#>   colMedians(matrix)          26.0         0.0860
#> colMedians(DelayedMatrix)      25.2         0.0932
#> colMedians(DelayedddgCMatrix) 850.0        107.0000
#>   colMedians(RleMatrix)       287.0         72.4000
#> colMedians(TricksyRleMatrix) 1550.0        333.0000
#>   colMedians(HDF5Matrix)       286.0         49.9000
#>   colMedians(matterMatrix)     139.0         40.5000
```

```
# Aside: apply(DelayedMatrix, 2, median) currently doesn't work
```

10 / 12

# Why DelayedMatrixStats?

## Backend-aware methods can improve performance

```
CS <- function(x, j) colSums(x[, j])          # DelayedArray
CS2 <- function(x, j) colSums2(x, cols = j)   # DelayedMatrixStats
j <- c(2001:3000, 5001:5500)
benchmark(CS(DelayedMatrix, j),              # Block-processing
          CS2(DelayedMatrix, j),              # Backend-aware
          CS(DelayeddgCMatrix, j),            # Block-processing
          CS2(DelayeddgCMatrix, j),           # Backend-aware
          CS(RleMatrix, j),                   # Block-processing
          CS2(RleMatrix, j),                  # Backend-aware
          times = 1)

#>               expr Median time (ms) Mem alloc (MB)
#>      CS(DelayedMatrix, j)          20.40         4.9200
#>     CS2(DelayedMatrix, j)           2.55         0.0244
#>      CS(DelayeddgCMatrix, j)       138.00        11.2000
#>     CS2(DelayeddgCMatrix, j)       20.30         1.4800
#>           CS(RleMatrix, j)         33.30        10.9000
#>          CS2(RleMatrix, j)         25.90         0.6650
```

# For more

**DelayedMatrixStats:** <https://github.com/PeteHaitch/DelayedMatrixStats>

**matter** developed by Kylie A. Bemis and available on Bioconductor

**matterArray:** <https://github.com/PeteHaitch/matterArray>

**Slides:** <http://peterhickey.org/presentations/>

**GitHub & Twitter:** @PeteHaitch